

Lecture 2: From MDP Planning to RL Basics

CS234: RL

Emma Brunskill

Spring 2017

Recap: Value Iteration (VI)

1. Initialize $V_0(s_i)=0$ for all states s_i ,
2. Set $k=1$
3. Loop until [finite horizon, convergence]
 - For each state s ,

$$V_{k+1}(s) = \max_a \left[r(s, a) + \gamma \sum_{s' \in S} p(s'|a, s) V_k(s') \right]$$

4. Extract Policy



Bellman
backup

V_k is optimal value if horizon=k

1. Initialize $V_0(s_i)=0$ for all states s_i ,
2. Set $k=1$
3. Loop until [finite horizon, convergence]
 - For each state s ,

$$V_{k+1}(s) = \max_a \left[r(s, a) + \gamma \sum_{s' \in S} p(s'|a, s) V_k(s') \right]$$

4. Extract Policy

Bellman
backup



Value vs Policy Iteration

- Value iteration:
 - Compute optimal value if horizon= k
 - Note this can be used to compute optimal policy if horizon = k
 - Increment k
- Policy iteration:
 - Compute infinite horizon value of a policy
 - Use to select another (better) policy
 - Closely related to a very popular method in RL: policy gradient

Policy Iteration (PI)

1. $i=0$; Initialize $\pi_0(s)$ randomly for all states s
2. Converged = 0;
3. While $i \neq 0$ or $|\pi_i - \pi_{i-1}| > 0$
 - $i=i+1$ V^π
 - Policy evaluation
 - Policy improvement

Policy Evaluation

1. Use minor variant of value iteration

$$V_{k+1}(s) = \underset{a}{\text{argmax}} \left[r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V_k(s') \right]$$

Policy Evaluation

1. Use minor variant of value iteration

$$V_{k+1}(s) = \max_a \left[r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V_k(s') \right]$$

$$V_{k+1}^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in S} p(s'|s, \pi(s)) V_k^\pi(s')$$

→ restricts action to be one chosen by policy

Policy Evaluation

1. Use minor variant of value iteration

$$V_{k+1}(s) = \max_a \left[r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V_k(s') \right]$$

$$V_{k+1}^\pi(s) = r(s, \pi(s)) + \gamma \sum_{s' \in S} p(s'|s, \pi(s)) V_k^\pi(s')$$

2. Analytic solution (for discrete set of states)

- Set of linear equations (no max!)
- Can write as matrices and solve directly for V

Policy Evaluation: Example

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Deterministic actions of TryLeft or TryRight
- Reward: +1 in state S1, +10 in state S7, 0 otherwise
- Let $\pi_0(s)$ =TryLeft for all states (e.g. always go left)
- Assume $\gamma=0$. What is the value of this policy in each s ?

$$V_{k+1}^{\pi}(s) = r(s, \pi(s)) + \gamma \sum_{s' \in S} p(s'|s, \pi(s)) V_k^{\pi}(s')$$

Policy Improvement

- Have $V^\pi(s)$ for all s (from policy evaluation step!)
- Want to try to find a better (higher value) policy
- Idea:
 - Find the state-action Q value of doing an action followed by following π forever, for each state
 - Then take argmax of Q_s

Policy Improvement

- Compute Q value of different 1st action and then following π_i

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^{\pi_i}(s')$$

- Use to extract a new policy

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

Delving Deeper Into Improvement

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^{\pi_i}(s')$$

$$\max_a Q^{\pi_i}(s, a) \geq V^{\pi_i}(s)$$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

- So if take $\pi_{i+1}(s)$ then followed π_i forever,
 - expected sum of rewards would be at least as good as if we had always followed π_i
- But new proposed policy is to always follow $\pi_{i+1} \dots$

Monotonic Improvement in Policy

- For any two value functions $V1$ and $V2$, let $V1 \geq V2 \rightarrow$ for all states s , $V1(s) \geq V2(s)$
- Proposition: $V^{\pi'} \geq V^{\pi}$ with strict inequality if π is suboptimal (where π' is the new policy we get from doing policy improvement)

Proof

$$\begin{aligned} V^\pi(s) &\leq \max_a Q^\pi(s, a) \\ &= \sum_{s' \in S} p(s' | s, \pi'(s)) \left[R(s, \pi'(s), s') + \gamma V^\pi(s') \right] \\ &\leq \sum_{s' \in S} p(s' | s, \pi'(s)) \left[R(s, \pi'(s), s') + \gamma \max_{a'} Q^\pi(s', a') \right] \\ &= \sum_{s' \in S} p(s' | s, \pi'(s)) \left[R(s, \pi'(s), s') + \right. \\ &\quad \left. \gamma \sum_{s'' \in S} p(s'' | s', \pi'(s')) (R(s', \pi'(s'), s'') + \gamma V^\pi(s'')) \right] \\ &\dots \leq V^{\pi'}(s) \end{aligned}$$

If Policy Doesn't Change ($\pi_{i+1}(s) = \pi_i(s)$ for all s) Can It Ever Change Again in More Iterations?

- Recall policy improvement step

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^{\pi_i}(s')$$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

Policy Iteration (PI)

1. $i=0$; Initialize $\pi_0(s)$ randomly for all states s
2. Converged = 0;
3. While $i \neq 0$ or $|\pi_i - \pi_{i-1}| > 0$
 - $i=i+1$
 - Policy **evaluation**: Compute V^π
 - Policy **improvement**:

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V^{\pi_i}(s')$$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

Policy Iteration Can Take At Most $|A|^{|S|}$ Iterations (Size of # Policies)

1. $i=0$; Initialize $\pi_0(s)$ randomly for all states s
2. Converged = 0;
3. While $i \neq 0$ or $|\pi_i - \pi_{i-1}| > 0$
 - $i=i+1$
 - Policy **evaluation**: Compute V^π
 - Policy **improvement**:

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V^{\pi_i}(s')$$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

* For finite state and action spaces

Policy Iteration

Fewer Iterations

More expensive per iteration

Value Iteration

More iterations

Cheaper per iteration


MDPs: What You Should Know

- Definition
- How to define for a problem
- MDP Planning: Value iteration and policy iteration
 - How to implement
 - Convergence guarantees
 - Computational complexity

Reasoning Under Uncertainty

Learn model of outcomes	Multi-armed bandits	Reinforcement Learning
Given model of stochastic outcomes	Decision theory	Markov Decision Processes
	Actions Don't Change State of the World	Actions Change State of the World

Reinforcement Learning



Goal: Maximize expected sum of future rewards

MDP Planning vs Reinforcement Learning

- No world models (or simulators)
- Have to learn how world works by trying things out

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

Policy Evaluation While Learning

- Before figuring out how should act
- 1st figure out how good a particular policy is (passive RL)

Passive RL

1. Estimate a model (and use to do policy evaluation)
2. Q-learning

Learn a Model

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Start in state S3, take TryLeft, go to S2
- In state S2, take TryLeft, go to S2
- In state S2, take TryLeft, go to S1
- What's an estimate of $p(s'=S2 \mid S=S2, a=\text{TryLeft})$?

Use Maximum Likelihood Estimate

E.g. Count & Normalize

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Start in state S3, take TryLeft, go to S2
- In state S2, take TryLeft, go to S2
- In state S2, take TryLeft, go to S1
- What's an estimate of $p(s'=S2 \mid S=S2, a=\text{TryLeft})$?
 - 1/2

Model-Based Passive Reinforcement Learning

- Follow policy π
- Estimate MDP model parameters from data
 - If finite set of states and actions: count & average
- Use estimated MDP to do policy evaluation of π

Model-Based Passive Reinforcement Learning

- Follow policy π
- Estimate MDP model parameters from data
 - If finite set of states and actions: count & average
- Use estimated MDP to do policy evaluation of π
- Does this give us dynamics model parameter estimates for all actions?
- How good is the model parameter estimates?
- What about the resulting policy value estimate?

Model-Based Passive Reinforcement Learning

- Follow policy π
 - Estimate MDP model parameters from data
 - If finite set of states and actions: count & average
 - Use estimated MDP to do policy evaluation of π
-
- Does this give us dynamics model parameter estimates for all actions?
 - No. But all ones need to estimate the value of the policy.
 - How good is the model parameter estimates?
 - Depends on amount of data we have
 - What about the resulting policy value estimate?
 - Depends on quality of model parameters


Good Estimate if Use 2 Data Points?

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Start in state S3, take TryLeft, go to S2, $r=0$
- In state S2, take TryLeft, go to S2, $r = 0$
- In state S2, take TryLeft, go to S1,
- What's an estimate of $p(s'=S2 \mid S=S2, a=\text{TryLeft})$?
 - $1/2$


Model-based Passive RL:

Agent has an estimated model in its head



Model-free Passive RL:

Only maintain estimate of Q



Q-values

- Recall that $Q^\pi(s,a)$ values are
 - expected discounted sum of rewards over H step horizon
 - if start with action a and follow π
- So how could we directly estimate this?

Q-values

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^{\pi_i}(s')$$

- Want to approximate the above with data
- Note if only following π , only get data for $a=\pi(s)$

Q-values

$$Q^{\pi_i}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V^{\pi_i}(s')$$

- Want to approximate the above with data
- Note if only following π , only get data for $a=\pi(s)$
- TD-learning
 - Approximate expectation with samples
 - Approximate future reward with estimate

Temporal Difference Learning


$$V^{\pi}(s) = r(s, \pi(s)) + \gamma \sum_{s' \in S} p(s'|s, \pi(s)) V^{\pi}(s')$$

- Maintain estimate of $V^{\pi}(s)$ for all states
 - Update $V^{\pi}(s)$ each time after each transition (s, a, s', r)
 - Likely outcomes s' will contribute updates more often
 - Approximating expectation over next state with samples
 - Running average

$$V_{samp}(s) = r + \gamma V^{\pi}(s')$$

$$V^{\pi}(s) = (1 - \alpha) V^{\pi}(s) + \alpha V_{samp}(s)$$

Decrease
learning rate
over time
(why?)



$$V_{samp}(s) = r + \gamma V^{\pi}(s')$$

$$V^{\pi}(s) = (1 - \alpha)V^{\pi}(s) + \alpha V_{samp}(s)$$

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Policy: TryLeft in all states, use $\alpha = 0.5$, $\gamma = 1$
- Set $V^{\pi} = [0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$,
- Start in state S3, take TryLeft, get $r=0$, go to S2
 - $V_{smp}(S3) = 0 + 1 * 0 = 0$
 - $V^{\pi}(S3) = (1-0.5)*0 + .5*0 = 0$ (no change!)

$$V_{samp}(s) = r + \gamma V^{\pi}(s')$$

$$V^{\pi}(s) = (1 - \alpha)V^{\pi}(s) + \alpha V_{samp}(s)$$

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Policy: TryLeft in all states, use alpha = 0.5, $\gamma=1$
- Set $V^{\pi}=[0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$,
- Start in state S3, take TryLeft, go to S2, get $r=0$
- $V^{\pi}=[0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$
- In state S2, take TryLeft, get $r=0$, go to S1
 - $V_{smp}(S2) = 0 + 1 * 0 = 0$
 - $V^{\pi}(S2)=(1-0.5)*0 + .5*0 = 0$ (no change!)

$$V_{samp}(s) = r + \gamma V^{\pi}(s')$$

$$V^{\pi}(s) = (1 - \alpha)V^{\pi}(s) + \alpha V_{samp}(s)$$

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Policy: TryLeft in all states, use $\alpha = 0.5$, $\gamma = 1$
- Start in state S3, take TryLeft, go to S2, get $r=0$
- In state S2, take TryLeft, go to S1, get $r=0$
- $V^{\pi} = [0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$
- In state S1, take TryLeft, go to S1, get $r=+1$
 - $V_{smp}(S1) = 1 + 1 * 0 = 1$
 - $V^{\pi}(S1) = (1-0.5)*0 + .5*1 = 0.5$

$$V_{samp}(s) = r + \gamma V^{\pi}(s')$$

$$V^{\pi}(s) = (1 - \alpha)V^{\pi}(s) + \alpha V_{samp}(s)$$

S1	S2	S3	S4	S5	S6	S7
Okay Field Site +1						Fantastic Field Site +10

- Policy: TryLeft in all states, use alpha = 0.5, $\gamma=1$
- Start in state S3, take TryLeft, go to S2, get $r=0$
- In state S2, take TryLeft, go to S1, get $r=0$
- $V^{\pi}=[0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$
- In state S1, take TryLeft, go to S1, get $r=+1$
- $V^{\pi}=[0.5 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$

Problems with Passive Learning

- Want to make good decisions
- Initial policy may be poor -- don't know what to pick
- And getting only experience for that policy

Can We Learn Optimal Values & Policy?

- Consider acting randomly in the world
- Can such experience allow the agent to learn the optimal values and policy?

Recall Model-Based Passive Reinforcement Learning

- **Follow policy π**
- Estimate MDP model params from observed transitions & rewards
 - If finite set of states and actions, count & avg counts
- Use estimated MDP to do policy evaluation of π

Recall Model-Based Passive Reinforcement Learning

- **Choose actions randomly**
- Estimate MDP model params from observed transitions & rewards
 - If finite set of states and actions, count & avg counts
- **Use estimated MDP to compute estimate of optimal value and policy**
- **Will policy converge to optimal value & policy**
 - (In limit of infinite data)?

Yes, if have reachability

- When acting randomly forever, still need to be able to visit each state and take each action many times
- Want all states to be reachable from any other state
- Quite mild assumption but doesn't always hold




Image source:
<http://ancient-heritage.blogspot.com/2014/05/crash-course-on-flying-in-face-of-logic.html>

Model-Free Learning w/Random Actions

- TD learning for policy evaluation:
 - As act in the world go through $(s, a, r, s', a', r', \dots)$
 - Update V^π estimates at each step
- Over time updates mimic Bellman updates
- Now do for Q values

Q-Learning

- Update $Q(s,a)$ every time experience (s,a,s',r)
 - Create new sample estimate

$$\begin{aligned}Q_{samp}(s, a) &= r + \gamma V(s') \\ &= r + \gamma \max_{a'} Q(s', a')\end{aligned}$$

- Update estimate of $Q(s,a)$

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha Q_{samp}(s, a)$$

Q-Learning Properties


- If acting randomly*, Q-learning converges Q^*
 - Optimal Q values
 - Finds optimal policy
- Off-policy learning
 - Can act in one way
 - But learning values of another π (the optimal one!)

*Again, under mild reachability assumptions

Towards Gathering High Reward

- Fortunately, acting randomly is sufficient, but not necessary, to learn the optimal values and policy
- Ultimately want to learn to get large reward

To Explore or Exploit?



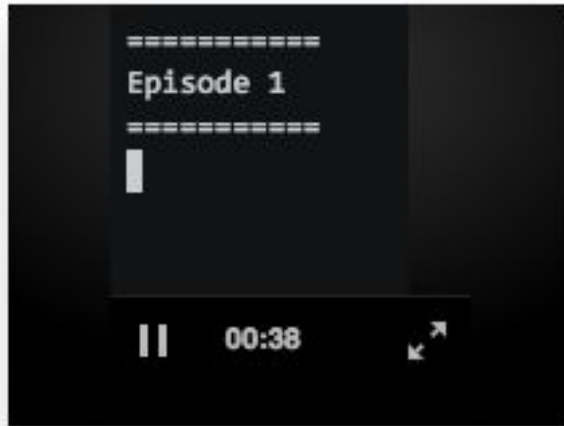
Simple Approach: E-greedy

- With probability $1-e$
 - Choose $\operatorname{argmax}_a Q(s,a)$
- With probability e
 - Select random action
- Guaranteed to compute optimal policy
- But even after millions of steps still won't always be following argmax of $Q(s,a)$

Greedy in Limit of Infinite Exploration (GLIE)

- E-Greedy approach
- But decay epsilon over time
- Eventually will be following optimal policy almost all the time
- We'll talk more about exploration/exploitation later in the course

Homework 1 Will Be Released This Week



FrozenLake-v0

Find a safe path across a grid of ice and water tiles.



FrozenLake8x8-v0

- Review/practice basic MDP planning
- Get familiar with Open AI gym for basic RL

What You Should Know

- Define MDP, Bellman operator, contraction, model, Q-value, policy
- Contrast MDP planning and RL
- Be able to implement
 - Value iteration, policy iteration, Q-learning and model-based RL
- Contrast benefits and weaknesses of Q-learning and model-based RL
 - On homework!
 - Data efficiency, computational complexity, etc.